

# **ML for Smarter Credit Card Fraud Detection**

## **From Raw Transactions to Insights**

ISYE 6740: Team 11

Wenxin Jiang, Yu Zheng, Yihan Liao, Qihao Cheng

# Motivation & Objective

We live in a world where billions of credit card transactions occur daily. With this massive volume of exchanges comes a critical challenge: **fraud detection**.

We aim to address the challenge of identifying fraudulent transactions within large-scale payment data using machine learning techniques. Our objectives are to **identify the best-performing model** and **enhance the precision of fraud detection** while maintaining efficiency in high-throughput environments.

| True                | False               |
|---------------------|---------------------|
| True Positive (TP)  | False Positive (FP) |
| False Negative (FN) | True Negative (TN)  |

Equations:

$$\text{False positive rate (FPR)} = \frac{\text{FP}}{\text{FP} + \text{TN}}$$

$$\text{False negative rate (FNR)} = \frac{\text{FN}}{\text{FN} + \text{TP}}$$

$$\text{Sensitivity} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

$$\text{Specificity} = \frac{\text{TN}}{\text{TN} + \text{FP}}$$

$$\text{Youden index} = \text{Sensitivity} + \text{Specificity} - 1$$

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

# Content

1. **Data Collection**
2. **EDA**
3. **Data Preparation**
4. **Methodology**
  - ☐ Logistic-Regression
  - ☐ SVM
  - ☐ Ensemble Method
  - ☐ Neural Network
5. **Discussion**
6. **Conclusion**

# Data Collection

- Dataset Source: Kaggle competition 'IEEE-CIS Fraud Detection', a collaboration between IEEE-CIS and Vesta Corporation.
- Dataset Size & Format: Large dataset derived from real-world e-commerce transactions, including features like device type and product characteristics.
- Scope: Improve fraud detection accuracy; reduce false positives to optimize fraud losses.
- Challenges: Heavy class imbalance typical in fraud detection keeping high false positives.



IEEE COMPUTATIONAL INTELLIGENCE SOCIETY · RESEARCH PREDICTION COMPETITION ·

## IEEE-CIS Fraud Detection

Can you detect fraud from customer transactions?

# EDA

# Data Overview

**432** features 【40(identity)+392(transaction)】

- **403** features after dropping high missing value column
- **2278** features after encoding.

**1** outcome variable [ “isFraud” ]

- 569877 Negative
- 20663 Positive

# Missing Value

After Review the percentage of missing value of each features.

Null Distribution are roughly shown as below:

100 %: 21 (features)

96 - 97%: 9

68 - 80%: 6

40 - 60%: 151

2 - 18%: 133


0 - 1%: 92

0: 22 (including outcome variable)

# Imbalance Data

Imbalance:

- Train set fraud ratio: 7.48%
- Validation set fraud ratio: 9.33%
- After Standardization, applied SMOTE to deal with the class imbalance by oversampling





# Data Preparation Overview

Step 1

- ✦✦ Missing Value Handling
  - Drop columns
  - Numerical
  - Categorical

Step 2

- ✦✦ Feature Scaling:
  - Standardization of continuous variables

Step 3

- ✦✦ Class Imbalance Solution:
  - SMOTE-NC

Step 4

- ✦✦ Feature Encoding:
  - One-hot encoding

# Step 1: Missing Value Handling

- Columns with >80% missing values dropped to prevent excessive noise.
- Numerical Features:
  - Before: Skewed distributions with NaNs.
  - After: Median imputation smoothened the distributions.
- Categorical Features:
  - Before: NaNs leading to incomplete categories.
  - After: Mode imputation preserved feature consistency.

# Step 3: Addressing Class Imbalance


Method: SMOTE-NC (Synthetic Minority Over-sampling Technique for Nominal and Continuous features)

Synthetic Sample Generation:

$$\mathbf{x}' = \mathbf{x} + \lambda(\mathbf{x}_{nn} - \mathbf{x})$$

Handling Categorical Features:

- Instead of interpolation, SMOTE-NC selects the **most frequent (mode)** value of the nominal feature among the **k-nearest neighbors**.



# Methodology

## Linear Models

- Logistic Regression

## Ensemble Methods

- Bagging
  - Random Forest
- Boosting
  - LightGBM,
  - XGBoost,
  - CatBoost

## Kernel/Nonlinear

- SVM
  - RBF SVM
  - Polynomial SVM
- Neural Networks

# Logistic Regression




Why Choose Logistic Regression?



Why the simplest?

- Interpretability
- Computational Efficiency
- Regulatory Issue
- Binary Classification



Assumptions: Likelihood maximization

- $P(Y|X)$





Tuning


- Regularization
- Encoding

# Logistic Regression - Result

- Excellent ROC (0.917)
- Mid PR (0.61 - 0.70 after categorical SMOTE)




# Ensemble Methods



# Why Bagging?

- Reduce overfitting
- Reduce the impact of outliers or noisy data points





# Random Forest - Methodology

1

## Create Bootstrap Samples

Randomly sample data points with replacement (bootstrap)

2

## Train Decision Tree


Each Tree use a subset of features and bootstrap samples. (sprt, 0.5). Tree grow independently.

3

## Aggregate Prediction

Majority Voting

# Random Forest - cont.



Why does it work well?

- “Teams” of Decision Trees



Tuning

```
(criterion='entropy', max_features='sqrt',  
max_samples=0.5, min_samples_split=80)
```



Assumptions


- Non Parametric
- Works well with mixed features
- Provide Feature Importance



Downsides


- “Black-box” Characteristics
- Trade-off between Data integrity and Computational force

# Random Forest - Result



- Excellent ROC (0.947)
- Mid PR (0.788 w. SMOTE)

# Important Features cont.



# Why Boosting?


## Advantages:

- **Handles Complex Relationships**
- **Improves Prediction Accuracy**
- **Reduces Bias and Variance**
- **Adaptable to Different Data Types**
- **Scalable and Efficient**

# XGBoost

(EXtreme Gradient Boost)

XGBoost:



## 1. Handling of Categorical Data

This is achieved through **partition-based splits**, where the algorithm determines the optimal way to divide categories into subsets to maximize **information gain**.

This approach reduces the need for extensive preprocessing and can lead to better performance by preserving the inherent relationships within categorical features.

## 2. Tree Splitting Method

XGBoost constructs trees using a **level-wise** growth strategy, meaning it expands the tree level by level.

At each level, it evaluates all possible splits across all leaves and selects the best splits to add to the tree.

The level-wise method in XGBoost helps maintain balanced trees, which can be beneficial for certain datasets and helps in controlling overfitting.


# LightGBM

Light Gradient Boosting Machine

## 1. Handling of Categorical Data

LightGBM can handle categorical features directly. By specifying which features are categorical, LightGBM applies a specialized algorithm to find the optimal split points for these features, often leading to better performance compared to traditional **one-hot encoding methods**.

LightGBM:



## 2. Tree Splitting Method

In contrast to XGBoost, LightGBM uses a **leaf-wise** growth strategy, splitting the leaf with the maximum loss reduction (gain).

This approach can lead to **deeper trees and better accuracy**.

However, it may also result in overfitting on smaller datasets, so parameters like `max_depth` are used to control tree depth.


# CatBoost

## Categorical Boosting

### 1. Handling of Categorical Data

CatBoost excels in processing categorical features without requiring extensive preprocessing. It employs a technique called **Ordered Target Statistics** to convert categorical data into numerical values during training.

This method calculates statistics (such as mean target values) for each category in a way that prevents data leakage and reduces overfitting. By integrating this approach within the training process, it maintains the **natural ordering and relationships of categorical features, leading to improved model accuracy.**



### 2. Tree Splitting Method

CatBoost employs a unique tree-building strategy known as **Symmetric Trees**. In this approach, all nodes at the same depth level in the tree are split using the same feature and threshold.

This symmetry simplifies the model and leads to faster predictions, as the structure of the tree is more regular.

Additionally, symmetric trees help in reducing overfitting and improve the model's generalization capabilities.



# XGBoost

Loss function

$$\mathcal{L}(\phi) = \sum_i l(\hat{y}_i, y_i) + \sum_k \Omega(f_k)$$

$$\text{where } \Omega(f) = \gamma T + \frac{1}{2} \lambda \|w\|^2$$

Pruning

Ridge Regression

$$\mathcal{L}^{(t)} = \sum_{i=1}^n l(y_i, \hat{y}_i^{(t-1)} + f_t(\mathbf{x}_i)) + \Omega(f_t)$$

Second order  
Taylor expansion

$$\mathcal{L}^{(t)} \simeq \sum_{i=1}^n [l(y_i, \hat{y}_i^{(t-1)}) + g_i f_t(\mathbf{x}_i) + \frac{1}{2} h_i f_t^2(\mathbf{x}_i)] + \Omega(f_t)$$


$$w_j^* = - \frac{\sum_{i \in I_j} g_i}{\sum_{i \in I_j} h_i + \lambda}$$

Output Value of leaf  $j$

$$\tilde{\mathcal{L}}^{(t)}(q) = - \frac{1}{2} \sum_{j=1}^T \frac{(\sum_{i \in I_j} g_i)^2}{\sum_{i \in I_j} h_i + \lambda} + \gamma T.$$

Scores of tree  $q$  (Impurity score)


# XGBoost



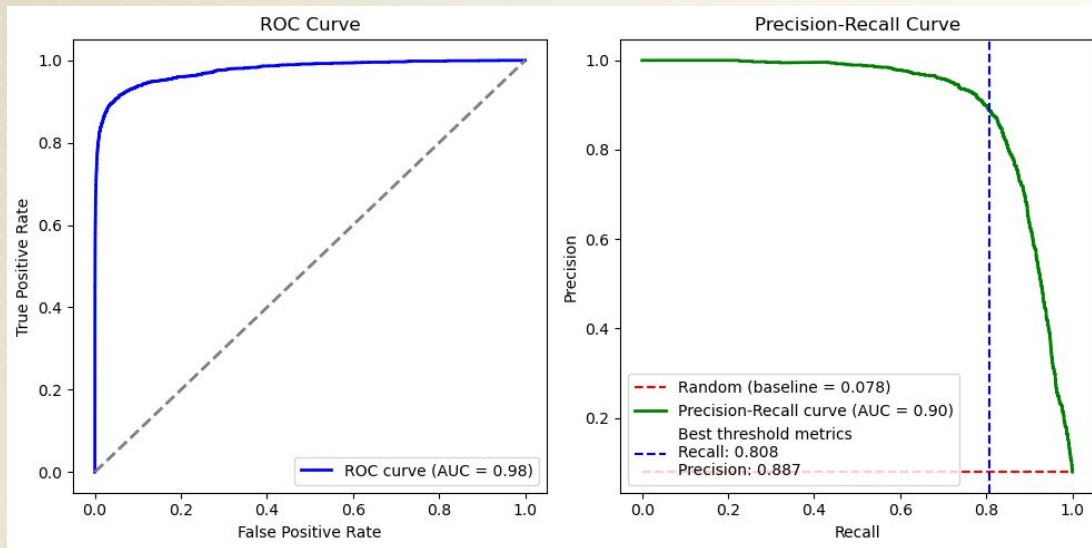
|                                   | Accuracy            | Best F1-score       | AUC PR             | AUC ROC            | Running Time        |
|-----------------------------------|---------------------|---------------------|--------------------|--------------------|---------------------|
| <b><u>1.Preprocessed Data</u></b> | <b><u>0.965</u></b> | <b><u>0.763</u></b> | <b><u>0.83</u></b> | <b><u>0.96</u></b> | <b><u>19.1s</u></b> |
| 2.Oversampled Data                | 0.964               | 0.752               | 0.82               | 0.95               | 33.5s               |

# XGBoost

Top 20 Most Important Features (XGBoost)




# LightGBM



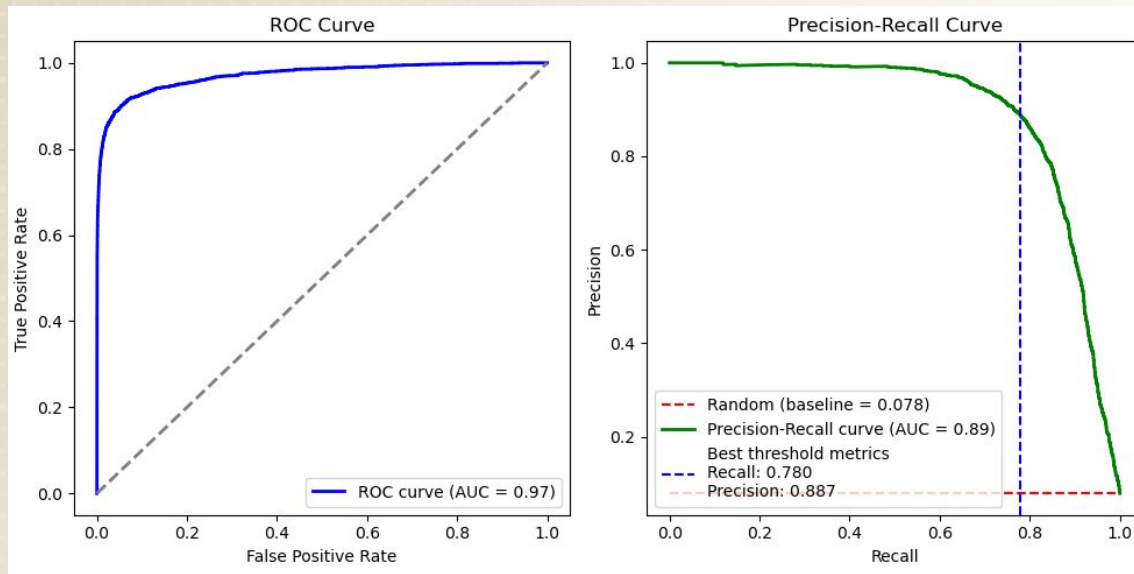
|                                   | Accuracy            | Best F1-score       | AUC PR             | AUC ROC            | Running Time        |
|-----------------------------------|---------------------|---------------------|--------------------|--------------------|---------------------|
| <b><u>1.Preprocessed Data</u></b> | <b><u>0.977</u></b> | <b><u>0.842</u></b> | <b><u>0.90</u></b> | <b><u>0.98</u></b> | <b><u>11.3s</u></b> |
| 2.Oversampled Data                | 0.973               | 0.820               | 0.88               | 0.97               | 29.6s               |

# LightGBM

Top 20 Most Important Features (LightGBM)




# CatBoost



|                                   | Accuracy            | Best F1-score       | AUC PR             | AUC ROC             | Running Time        |
|-----------------------------------|---------------------|---------------------|--------------------|---------------------|---------------------|
| <b><u>1.Preprocessed Data</u></b> | <b><u>0.974</u></b> | <b><u>0.830</u></b> | <b><u>0.89</u></b> | <b><u>0.972</u></b> | <b><u>55.2s</u></b> |
| 2.Oversampled Data                | 0.971               | 0.813               | 0.87               | 0.968               | 97.6s               |

# CatBoost

Top 20 Most Important Features (CatBoost)





# Experimental Results

1. LightGBM performs best.
2. Oversample does **not** provide helpful information for boosting model training.

|                        | Accuracy            | Best F1-score       | AUC PR             | AUC ROC            | Running Time        |
|------------------------|---------------------|---------------------|--------------------|--------------------|---------------------|
| XGBoost                | 0.965               | 0.763               | 0.83               | 0.96               | 19.1s               |
| <b><u>LightGBM</u></b> | <b><u>0.976</u></b> | <b><u>0.842</u></b> | <b><u>0.90</u></b> | <b><u>0.98</u></b> | <b><u>11.3s</u></b> |
| CatBoost               | 0.974               | 0.830               | 0.89               | 0.97               | 55.2s               |

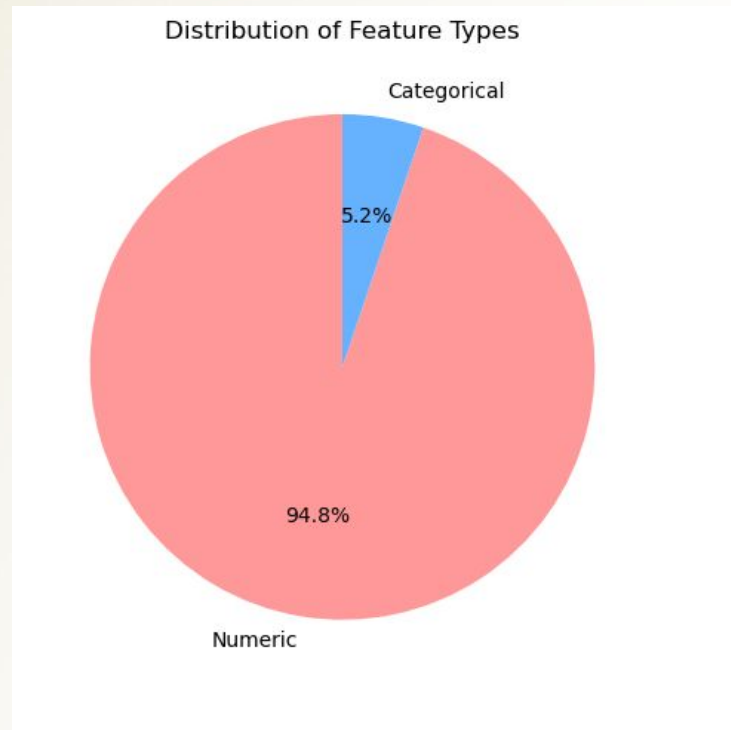
# Discussion

- Categorical features are important in classifying the fraud transaction  
The number of categorical features in top 20 most important features (245): Random Forest (2), XGBoost (1), LightGBM(5), CatBoost(5)
- The original dataset is imbalanced, but due to the large number of samples, boosting models can effectively learn the relationship between transactions and fraud. However, oversampling introduces noise into the data, which boosting models tend to fit, ultimately compromising their performance.

# Discussion

## Low Proportion of Categorical Features


- The dataset contains **few categorical features (less than 10%)** and is dominated by numerical features, so LightGBM's efficiency in handling numerical data gives it an edge.
- CatBoost's key advantage lies in its advanced handling of categorical features, so when categorical features are **sparse**, this advantage becomes less impactful.



# Discussion

## Category Cardinality

- If the categorical features have **low cardinality**, LightGBM's processing can yield comparable results to CatBoost without requiring its specialized encoding.





# Support Vector Machines(SVMs)

- Radial Basis Function(RBF) SVM
- Polynomial Kernel SVM

# RBF Kernel SVM

- Training Model
  - Selected 20% data from train dataset as sample, then PCA was applied to reduce dimensionality while retaining 90% variance.
  - The same process applied to two data sets
    - Preprocessed Data
    - Oversampled Data
  - Grid Search, best model setting  $C = 100$ ,  $\gamma = 0.001$

# RBF Kernel SVM



|                                   | Accuracy             | F1-score             | AUC PR               | AUC ROC              | Running Time         |
|-----------------------------------|----------------------|----------------------|----------------------|----------------------|----------------------|
| <b><u>1.Preprocessed Data</u></b> | <b><u>0.9505</u></b> | <b><u>0.5908</u></b> | <b><u>0.6733</u></b> | <b><u>0.8886</u></b> | <b><u>2m 30s</u></b> |
| 2.Oversampled Data                | 0.8720               | 0.4857               | 0.6280               | 0.9016               | 15m                  |

# RBF Kernel SVM


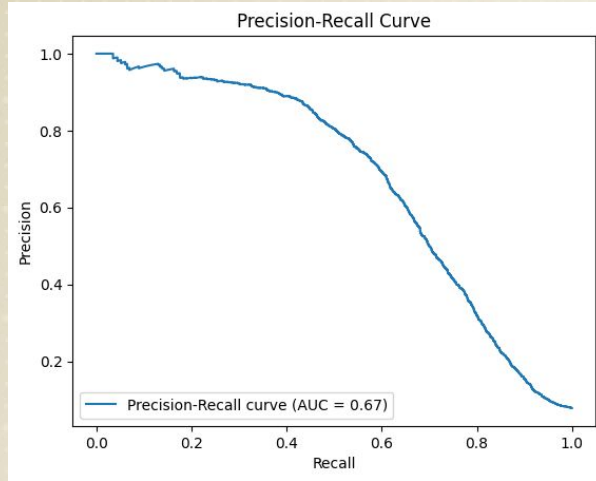
- Performance Evaluation
  - Preprocessed Data:
    - Higher overall accuracy and precision for fraud detection. Poor recall for fraud cases, missing more fraudulent transactions.
  - Oversampled Data:
    - Improved recall for fraud (77% vs. 46%) but at the cost of significantly lower precision. Lower F1 score for fraud indicates the trade-off between precision and recall is less favorable.
  - Both models face difficulties balancing precision and recall due to the highly imbalanced nature of the dataset.



# Polynomial Kernel SVM

- Training Model
  - Selected 20% data from train dataset as sample (34026, 2680), then PCA was applied to reduce dimensionality while retaining 90% variance.
  - The same process applied to two data sets
    - Preprocessed Data
    - Oversampled Data
  - Grid Search, best model setting  $C = 10$ , degree = 3, gamma = scale, coef0 = 1

# Polynomial Kernel SVM



|                                   | Accuracy             | F1-score             | AUC PR               | AUC ROC              | Running Time         |
|-----------------------------------|----------------------|----------------------|----------------------|----------------------|----------------------|
| <b><u>1.Preprocessed Data</u></b> | <b><u>0.9513</u></b> | <b><u>0.6303</u></b> | <b><u>0.6733</u></b> | <b><u>0.8861</u></b> | <b><u>1m 30s</u></b> |
| 2.Oversampled Data                | 0.8824               | 0.5137               | 0.6236               | 0.9110               | 10m 20s              |

# Polynomial Kernel SVM

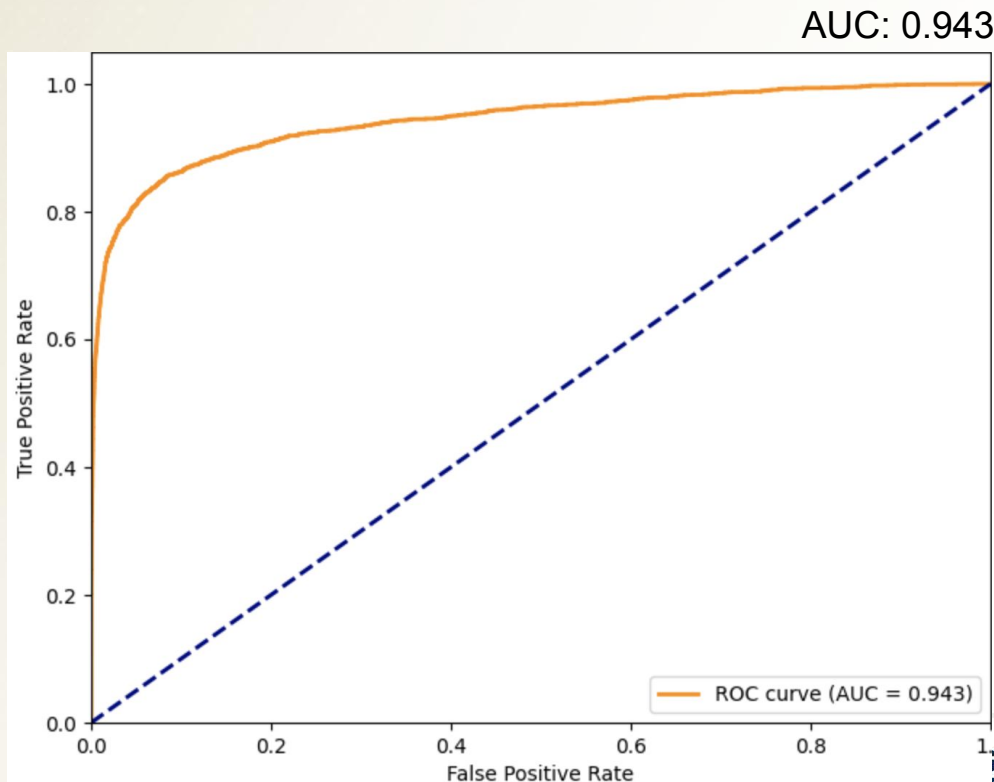
- Performance Evaluation
  - Preprocessed Data:
    - Higher overall accuracy, precision, and F1 score for fraud detection. Low recall (53%) indicates many fraud cases are missed.
  - Oversampled Data:
    - Improved recall (79%) for fraud but at the expense of significantly lower precision (38%). Lower accuracy (88.2%) and F1 score (0.5137) for fraud detection due to an increase in false positives.
  - Polynomial SVMs, especially with higher degrees, are prone to overfitting, particularly when the dataset contains synthetic features or noise (as introduced by SMOTE).

# Neural Networks

```
# Training configurations
class Config:
    # Model parameters
    model_architecture = {
        'hidden_layers': [128, 64, 32, 16],
        'dropout_rate': 0.01,
    }

    # Training parameters
    num_epochs = 50
    steps_per_epoch = 200
    batch_size = 32
    learning_rate = 0.001
    fraud_ratio = 0.5 # For balanced batches

    # Encoder parameters
    encoder_params = {
        'onehot_threshold': 10,
        'outlier_threshold': -999,
        'std_threshold': 5,
        'cache_dir': 'encoder_cache',
        'force_recompute': False
    }
```




# Neural Networks

- **Hidden Layers:** Configured with dimensions [128, 64, 32, 16], progressively reducing complexity while retaining the ability to learn deep representations.
- **Activation Function:** Each layer uses ReLU (Rectified Linear Unit) activation to introduce non-linearity, enabling the model to learn complex patterns.
- **Regularization:** A dropout layer with a rate of 0.01 is applied after each hidden layer to mitigate overfitting by randomly deactivating neurons during training.
- **Output Layer:** A single neuron with no activation directly outputs logits, suitable for binary classification tasks when paired with a sigmoid function during evaluation.

# Neural Networks

Training and Test Loss over Epochs



**Checkpoint saved at epoch 40**

True Positive Rate: 0.7961

True Negative Rate: 0.9435

False Positive Rate: 0.0565

False Negative Rate: 0.2039

Precision: 0.5454

F1 Score: 0.6473

# Neural Networks

The model is trained using binary **cross-entropy loss** with logits, which combines a sigmoid activation with cross-entropy for stable binary classification. **The Adam optimizer**, with a learning rate of 0.001, adapts learning rates for efficient convergence. Training is conducted over 50 epochs with 200 steps per epoch, ensuring adequate exposure to patterns in the dataset. **Balanced batch sampling**, with a fraud ratio of 50%, addresses class imbalance by including sufficient minority class instances during training. To enhance generalization, **dropout regularization** (rate 0.01) mitigates overfitting by randomly deactivating neurons.

The training process incorporates periodic model checkpointing for resumability and generates validation metrics and visualizations, such as AUC history and ROC curves, to monitor progress and ensure robust evaluation. These features make the model both effective and straightforward to extend or adapt for fraud detection tasks.



# Discussion

## Model Performance Comparison

- Model with High Complexity tends to capture the trend in this high dynamic dataset
  - The Threshold of a good model is set as 0.8 to ensure bearable “false alarm cost”

|         | Log Reg | SVM    | RFC  | Boost | NN   |
|---------|---------|--------|------|-------|------|
| ROC-AUC | .917    | 0.8886 | .947 | .980  | .943 |
| PR-AUC  | .700    | 0.6733 | .788 | .900  | .809 |



**Questions?  
&  
Thank you!**